

# EE 599 ML Systems Final Project

## Efficient Fine-tuning of LLM on a Single GPU

Instructor: Arash Saifhashemi  
TA: Lei Gao, Kevin Yang

### 1 Introduction

Transformer-based large language models (LLMs) like OpenAI's GPT-4 and ChatGPT have significantly expanded language generation and understanding capabilities. Trained on extensive web-scale datasets, these models excel not only in natural language processing but also in solving complex tasks through adept instruction handling and multi-step reasoning. This makes LLMs pivotal in advancing towards artificial general intelligence (AGI).

In our project, we focus on instruction tuning as a key method for adapting pretrained LLMs, using Meta AI's open-sourced LLaMA2 7B model. We'll investigate efficient fine-tuning techniques like gradient accumulation, checkpointing, mixed precision training, and low-rank adaptation, which are crucial for optimizing LLMs in resource-constrained environments, particularly for mitigating concerns regarding memory usage.

### 2 Background

In this section, your task is to review the LLM survey paper [?] and the LLaMA papers [?, ?] selectively, and provide comprehensive answers to the questions below. These questions will serve as a guide for reading and understanding the content of the papers. You may use Google or ChatGPT to assist in your research, but ensure that your answers are accurate and demonstrate a clear understanding of the concepts discussed in the papers.

1. What are the four major aspects of LLMs covered in the LLM survey paper?
2. What are the three major differences between LLMs and PLMs? What are the three typical emergent abilities for LLMs?
3. Where are pre-training data from?
4. What are the three main types of instruction tuning datasets? What is the Alpaca dataset [?]? Pick a training sample and describe it.
5. What are the pertaining data used by LLaMA? How many tokens are in the entire training set for training LLaMA?
6. Before pre-training, what is the procedure for data preprocessing?
7. What is tokenization? Name a few tokenization algorithms. What is the tokenization method used by LLaMA?
8. What are the main three types of transformer architecture? Name a few models for each type.
9. Why are LLMs mainly decoder-only architecture?
10. What is position embedding? Name a few position embedding algorithms.
11. What is language modeling? What is the difference between causal language modeling and masked language modeling?

12. How to generate text from LLMs [?]? Explain different decoding strategies.
13. What is instruction tuning, and why is it important?
14. What is alignment tuning, and why is it important?
15. What is parameter-efficient fine-tuning, and why is it important? What is the difference between prefix tuning and prompt tuning?
16. What is prompt engineering, and why is it important? What is zero-shot and few-shot demonstrations?
17. What is the difference between in-context learning and chain-of-thought prompting?
18. What are the three basic types of ability evaluation for LLMs? What is perplexity [?]? What is hallucination?
19. What is human alignment, and why is it important?
20. Name a few comprehensive evaluation benchmarks for the evaluation of LLMs.

Your answer:

1. Pre-training, adaption tuning, utilization and capacity evaluation.
2. Three differences between LLMs and PLMs: Emergent abilities, development and usage approach, research and engineering integration. Three typical emergent abilities for LLMs: In-context learning, instruction following, step-by-step reasoning.
3. Pre-training data for LLMs come from massive text corpora collected from internet. These corpora are typically collected from diverse sources across the internet, including websites, books, articles, and other digital text repositories. This wide-ranging collection helps ensure the models learn a broad understanding of language and context.
4. Three main types of instruction tuning datasets: Crowd sourced datasets, expert curated datasets, automatically generated datasets. Alpaca dataset is using as instruction tuning to enhance the performance of LLMs. One example task is the model to categorize a list of fruits into different classes based on given attributes. Which input is the list of fruits and the expected output is organize these fruits into different categories.
5. LLaMA models are trained on a large scale dataset, have been trained using 1.4 trillion tokens.
6. Before pre-training, first we need to gathering data and normalize it, splitting into pieces or tokens, removing the sensitive information and dividing the dataset into training, validation, test sets, finally formatting the data.
7. Tokenization is the process of converting text into smaller units. Example algorithms like word tokenization, subword tokenization and charactor tokenization. Tokenization used in LLaMA for variety of languages and scripts by subword tokenization method.
8. Main three types of transformer architecture are encoder-only transformers, decoder-only transformers and encoder-decoder transformers. Examples for encoder-only like, BERT which is designed for pre-trained deep bidirectional representations, DistilBERT which is a smaller and faster version of BERT. Examples for decoder-only like, GPT which is a generative pre-trained transformer. Example for encoder-decoder like, T5 which is a text-to text transfer transformer to frames all NLP tasks as text to text problem and BART which combines BERT and GPT to understand and the generate text.
9. One reason is because the decoder-only models have a simplified architecture which can reducing complexity and easier to train. Another reason is because decoder-only architectures can scaled up with more parameters which can generate more complex text output as we want.
10. Position embedding is a technique used to incorporate information about absolute position of tokens in the input. Example algorithms like leaned position embedding and relative position embedding.
11. Language modeling is a task in natural language processing to predicting the probability distribution of words. CLM uses a one-directional context to predict the next word in sequence while MLM leverages a bidirectional context to predict a word in sequence.
12. First choose the word with the highest probability at each step and then consider the multiple high probability word in sequence and finally introduce randomness into generation process to make outputs

more diverse. For greedy search is quick but can lead to repetitive text, for beam search can maintains high probability sequences and for random sampling can introduce randomness.

13. Instruction tuning is a method to refine the performance of language models. It is important because it can enable the model to understand and execute complex instructions to enhance its applicability across diverse real world applications.

14. Alignment tuning is a specialized approach to adapting language models to better align their responses with human intentions. It is important because it can enhancing the safety and societal acceptance of AI systems which can help the model behave responsible.

15. Parameter-efficient fine-tuning is a technique to keep the parameters of pre-trained model fixed. It is important because it allows large modifications with minimal computational resources, make it easier to adapt large models. Prefix tuning adds trainable parameters to the input while prompt tuning prepended to the input sequence to steer the responses of model to the desired outcome.

16. Prompt engineering is to guild the behavior of language models to produce desired outputs. It is important because it maximized the efficiency of models without extensive retraining. Zero-shot is the model performs a task without any prior examples and few-shot is the model given a few examples to demonstrate the task before asked.

17. In-context learning is how the model to understand and respond a task based on the provided context. Chain-of-thought prompting is the process of intermediate reasoning in prompt and guiding the model to follow the logical sequence to the conclusion.

18. Human alignment is the process of ensuring the AI system behave with human values, ethics and intentions. It is important because it reduce the risks of AI behaviors which might harmful for human.

19. Three basic types of ability evaluation for LLMs are performance specialization, generalization of abilities and behavioral tests. Perplexity is a metric that used to evaluate language models like GPT-2. Hallucination is the way to instance which model generates the text that is ungrounded to the input data.

20. Examples like GLUE, general language understanding evaluation which is a collection of nine different tasks designed to measure a model's ability to understand language. SQuAD, standford question answering dataset which focusing on the ability of answering questions based on give text.

## 3 Optimization Methods

### 3.1 Gradient Accumulation

In standard neural network training, we typically divide our data into smaller chunks called mini-batches and process them sequentially. The network generates predictions for each batch, and we compute the loss by comparing these predictions to the actual targets. Afterward, we perform a backward pass to calculate gradients, which are used to adjust the model's weights in the direction of these gradients.

Gradient accumulation alters this last step of the training process. Instead of updating the model's weights after processing each individual batch, we save the gradient values and continue to the next batch, accumulating the new gradients. Weight updates are deferred until the model has processed several batches.

The primary purpose of gradient accumulation is to simulate a larger effective batch size. For instance, imagine you intend to use a batch size of 32 images, but your hardware can handle only up to 8 images without running out of memory. In this scenario, you can work with smaller batches of 8 images and update the weights after every 4 batches by accumulating gradients from each batch in between. To ensure the accumulated gradient remains equivalent to the non-accumulated gradient on a batch of 32 images, you must divide your loss value on each mini-batch of 8 images by 4. Consider a linear model with MSE loss, and mathematically explain why this division step can yield identical results.

Your answer: Consider a simple linear model  $f(x) = wx + b$ , where  $w$  and  $b$  are the model's weight and bias, respectively. The Mean Squared Error (MSE) loss for a set of predictions  $\hat{y}$  and targets  $y$  is given by:

$$MSE = \frac{1}{n} \sum_{i=1}^n (\hat{y}_i - y_i)^2$$

where  $n$  is the number of samples in the batch. The gradient of the loss with respect to the weight  $w$  is

computed as follows:

$$\frac{\partial MSE}{\partial w} = \frac{2}{n} \sum_{i=1}^n (\hat{y}_i - y_i) \cdot \frac{\partial \hat{y}_i}{\partial w}$$

Given  $\hat{y}_i = wx_i + b$ , where  $\frac{\partial \hat{y}_i}{\partial w} = x_i$ , the gradient simplifies to:

$$\frac{\partial MSE}{\partial w} = \frac{2}{n} \sum_{i=1}^n (\hat{y}_i - y_i) x_i$$

Due to hardware limitations, consider using a batch size of 8 and accumulating gradients over 4 such mini-batches to simulate an effective batch size of 32. The gradient for each mini-batch of 8 samples is:

$$\frac{\partial MSE_{mini}}{\partial w} = \frac{2}{8} \sum_{i=1}^8 (\hat{y}_i - y_i) x_i$$

After processing 4 mini-batches, the accumulated gradient becomes:

$$Accumulatedgradient = 4 \times \frac{2}{8} \sum_{all32samples} (\hat{y}_i - y_i) x_i = \frac{2}{2} \sum_{all32samples} (\hat{y}_i - y_i) x_i$$

If we divide the loss of each mini-batch by 4, the gradient for each batch becomes:

$$\frac{\partial MSE_{mini}}{\partial w} = \frac{1}{4} \cdot \frac{2}{8} \sum_{i=1}^8 (\hat{y}_i - y_i) x_i$$

Thus, accumulating this gradient over 4 batches yields:

$$Accumulatedgradient = \frac{1}{4} \cdot 4 \times \frac{2}{8} \sum_{all32samples} (\hat{y}_i - y_i) x_i = \frac{2}{32} \sum_{all32samples} (\hat{y}_i - y_i) x_i$$

This adjustment ensures that the accumulated gradient matches the gradient computed for a single batch of 32 samples, thereby maintaining consistency in learning updates. This shows that dividing the MSE loss by the number of mini-batches (4 in this case) ensures the accumulated gradient matches the gradient computed for a single batch of 32 samples.

However, it's important to note that gradient accumulation may not always yield identical results, especially for deep neural network architectures that employ batch-wise regularization methods like batch normalization. Batch normalization relies on statistics calculated within each batch, and when you accumulate gradients, these statistics may not be consistent. Mathematically explain the batch normalization layer behavior and why gradient accumulation yields non-identical results in this case.

**Your answer:** Batch normalization involves normalizing the input layer by adjusting and scaling the activations. Mathematically, for a given feature in a batch, batch normalization is defined as:

$$BN(x_i) = \gamma \left( \frac{x_i - \mu_B}{\sqrt{\sigma_B^2 + \epsilon}} \right) + \beta$$

where  $x_i$  are the inputs to the layer,  $\mu_B$  is the mean of the batch,  $\sigma_B^2$  is the variance of the batch,  $\gamma$  and  $\beta$  are parameters to be learned, and  $\epsilon$  is a small constant added for numerical stability.

The key point in batch normalization is that it relies on the mean ( $\mu_B$ ) and variance ( $\sigma_B^2$ ) computed from the current batch. When using gradient accumulation, each mini-batch contributes to the accumulated gradient, but the batch normalization statistics ( $\mu_B$  and  $\sigma_B^2$ ) are recalculated for each mini-batch. This means the normalization is performed using different means and variances across these batches. Therefore, even though gradients are accumulated, the differing batch statistics lead to discrepancies in the normalized inputs each time, potentially affecting the learning dynamics. Consequently, the updates computed through gradient accumulation with batch normalization may not perfectly simulate the updates that would be computed in a single larger batch.

## 3.2 Gradient Checkpointing

Gradient checkpointing is another memory-saving technique that strikes a balance between memory and computation during neural network training. While it's well understood that storing all activations and intermediate results from the forward pass is essential for performing backward propagation, this process can lead to a significant increase in memory usage, often referred to as the "memory blow-up" problem. When monitoring memory consumption during model training, you'll notice that it gradually accumulates and reaches its peak after completing the forward pass. Explain why model inference does not have such "memory blow-up" problem.

Your answer: Model inference doesn't experience the "memory blow-up" problem seen in training primarily because it only involves the forward pass of the neural network. During inference, the network computes output predictions from given inputs, using the learned weights and biases. Crucially, it does not require the storage of intermediate activations needed for backpropagation, as no learning or adjustment of weights occurs. This means the memory footprint is much smaller during inference compared to training, where memory needs to accommodate both the activations for computing gradients and the gradients themselves.

Instead of retaining all activations and intermediate results in memory throughout the forward pass, gradient checkpointing offers an alternative approach. It strategically selects checkpoints within the network where memory efficiency takes precedence over computation time. When the algorithm encounters these checkpoints during the subsequent backward pass, it does not rely on stored activations because they were not stored in memory in the first place. Instead, it recalculates them by re-executing the forward pass up to the specific checkpoint.

To optimize memory usage and forward computation steps using gradient checkpointing in a feed-forward neural network composed of  $n$  layers, let's consider two strategies: Imagine a simple feed-forward neural network with  $n$  layers, each generating activations of the same size. Upon completing the forward pass, the algorithm retains all activations, resulting in an  $O(n)$  memory requirement for this network. In other words, it consumes  $n$  units of memory. In terms of computation, the forward pass requires  $O(n)$  steps, assuming that each layer requires 1 unit of computation.

Now, the question revolves around determining the optimal placement of checkpoints to strike a balance between memory and computation. A poor strategy would be discarding all activations during the forward pass and recomputing them later when needed. Conversely, a more general and optimal strategy suggests discarding activations every  $\sqrt{n}$  steps. What are the memory requirement and forward computation steps for the two strategies in big O notation? Check this medium post [?] for more illustrations.

Your answer: The two strategies in gradient checkpointing for managing memory and computation in a feed-forward neural network with  $n$  layers are summarized below in terms of their big O notation for memory and computation requirements: **Strategy 1 (Discarding all activations during forward pass):** The memory requirement is  $O(1)$ , as only the current activation is stored at any time, with all others discarded and recalculated as needed. The forward computation steps are  $O(n^2)$ , since every time a gradient is needed for a layer during the backward pass, the forward computations for all preceding layers up to that point must be redone, leading to a quadratic increase in computational steps. **Strategy 2 (Checkpointing every  $\sqrt{n}$  steps):** The memory requirement is  $O(\sqrt{n})$ , with checkpoints saved after every  $\sqrt{n}$  layers, resulting in approximately  $\sqrt{n}$  checkpoints stored. The forward computation steps are  $O(n\sqrt{n})$ , as for each of the  $n$  layers, if its activation needs to be recomputed during the backward pass, it involves re-executing the forward pass from the nearest preceding checkpoint, averaging out to  $\sqrt{n}$  recomputation steps per layer.

In practical scenarios, neural networks often feature layers with varying activation sizes, and the computational complexity of each layer's forward pass can differ significantly. Furthermore, the forward graph of these networks may not adhere to a strictly sequential pattern. Consequently, the task of strategically placing checkpoints to optimize memory usage and computation steps remains an active area of research and attention.

## 3.3 Low-Rank Adaptation (LoRA)

Parameter-efficient Fine-tuning (PEFT) is a technique in NLP that enhances the performance of pre-trained language models on specific tasks. It freezes most of the pre-trained model's parameters and only fine-tunes a subset of layers, reducing computational and memory demands and training time, while retaining a good

performance as possible. The idea of reducing the number of trainable parameters has its roots in transfer learning within computer vision tasks, where traditionally, only the final fully connected layer was updated.

Before we dive into the technical aspect of Low-Rank Adaptation (LoRA), one must revisit the transformer architecture and understand the mechanism of self-attention. In addition, LoRA uses the concept of truncated Singular Value Decomposition (SVD). Please study and understand the concepts of SVD and truncated SVD. Answer the following questions:

1. What is matrix rank?
2. What are three decomposed matrices by SVD?
3.  $U$  and  $V$  are orthogonal matrices. Why does it imply  $UU^T = I, VV^T = I$ ?
4. If a matrix  $W \in R^{n \times n}$  is full rank, what is its rank?
5. Suppose a full rank matrix  $W \in R^{n \times n}$  represents an image. After we apply SVD to this matrix, we modify the singular matrix by only keeping its top- $k$  singular values and discarding the rest (i.e., set the rest of the singular values to zero). Then, we reconstruct the image by multiplying  $U$ , modified  $S$ , and  $V$ . What would the reconstructed image look like? What if you increase the values of  $k$  (i.e., keep more singular values)?
6. If a matrix  $W \in R^{n \times n}$  is low rank, what does its singular matrix look like?
7. If the top- $k$  singular values of a matrix  $W \in R^{n \times n}$  are large, and the rest are near zero, this matrix  $W$  exhibits low-rank or near-low-rank behavior. Can you represent  $W$  by two low-rank matrices,  $A$  and  $B$ ? If so, what are those two matrices' expressions in terms of  $U$ ,  $S$ , and  $V$ ? Do you think those two matrices are a good approximation of  $W$  (i.e.,  $W \approx AB$ )?
8. The above operation is called truncated SVD. Under what situation do you think truncated SVD fails to make a good approximation? Think about the singular matrix.

Your answer:

1. The matrix rank is a measurement of dimension of the vector space spanned by its rows or columns. It tells how much information contained in the matrix and how many dimensions can be described by rows or columns.
2. Let  $A$  be a sample matrix. The first matrix is  $U$ , which is an  $m \times m$  orthogonal matrix. The columns of  $U$  are the eigenvectors of  $AA^T$ . The second matrix is  $\Sigma$ , which is an  $m \times n$  diagonal matrix containing the singular values of  $A$ . The third matrix is
3. As  $U$  is an orthogonal matrix, the columns of  $U$  are orthonormal. Thus, each element  $u_{ij}$  in matrix  $U$  is the dot product of column  $i$  of  $U$  with column  $j$  of  $U$ . Since the columns are orthonormal, the dot product of any column with itself should be 1 and the dot product of different columns should be 0. Therefore,  $UU^T = I$ . Similarly, for  $V^T$ , changes pertain to the dot product of row  $i$  of  $V^T$  with row  $j$  of  $V^T$ . The diagonal entries will be the dot product of each column with itself, which is 1, and the off-diagonal entries will be zero. Therefore,  $V^TV = I$ .
4. Full rank means  $W$  is an  $n \times n$  matrix and all  $n$  columns are linearly independent, so the rank is  $n$ .
5. Retaining only top- $k$  singular values in  $\Sigma$  means lowering the singular values, causing the reconstructed image to lose some details. If we keep more singular values, as  $k$  increases, it will improve the fidelity. The image will become sharper and has more details in it. Overall, smaller  $k$  compresses the image while larger  $k$  emphasizes fidelity to the original image.
6. Its singular matrix will have sparse non-zero entries which means most singular values will be zero. The singular values in  $\Sigma$  are arranged in descending order which means the largest singular values will appear first on the diagonal, followed by smaller values and then zeros.
7. First, set  $U_k$  as the new  $U$  containing the first  $k$  columns,  $\Sigma_k$  as the new  $\Sigma$  containing the top left  $k \times k$  blocks of  $\Sigma$  with the largest  $k$  singular values, and  $V_k$  as the new  $V$  containing the first  $k$  columns. Define  $A = U_k \Sigma_k^{1/2}$  and  $B = \Sigma_k^{1/2} V_k^T$ . Now we can compute  $AB = U_k \Sigma_k^{1/2} \Sigma_k^{1/2} V_k^T = U_k \Sigma_k V_k^T$ , which is the approximation to  $W$  in terms of the Frobenius norm. This is a good approximation if  $k$  is large enough.
8. It is a good approximation method. However, if the lower singular values contain a significant amount

of information, it could fail as there might be too much information relevant to the data structure. The omission can lead to substantial loss of information.

Now, you are equipped with the knowledge to understand LoRA. As one of the most effective PEFT techniques, proposed by Edward Hu et al. in 2021 [?], the success of LoRA is based on an important observation that both pre-trained weights and change of weights during fine-tuning for LLMs are low-rank matrices.

To illustrate this, consider representing the final optimal weights after fine-tuning as  $W^* = W_0 + \Delta W$ , where  $W_0 \in R^{n \times n}$  represents the pre-trained weights, and  $\Delta W \in R^{n \times n}$  captures the weight changes during fine-tuning. Notably,  $\Delta W$  can be approximated as the product of two matrices,  $AB$ , with  $A \in R^{n \times r}$  and  $B \in R^{r \times n}$ . This implies that during fine-tuning, we don't have to update the whole matrix  $W_0$ . Instead, we modify the forward pass in the layer, changing it from  $y = xW_0$  to  $y = xW_0 + xAB = x(W_0 + AB)$ , as depicted in Figure ?? . We freeze  $W_0$  and only update  $A$  and  $B$ . Notice that since  $r \ll n$ , the total number of parameters of  $A$  and  $B$  is much smaller than that of  $W_0$ . To make the modified forward pass identical to the original one at the beginning of fine-tuning,  $A$  is randomly initialized in Gaussian distribution, and  $B$  is initially set to 0, so  $AB = 0$ .

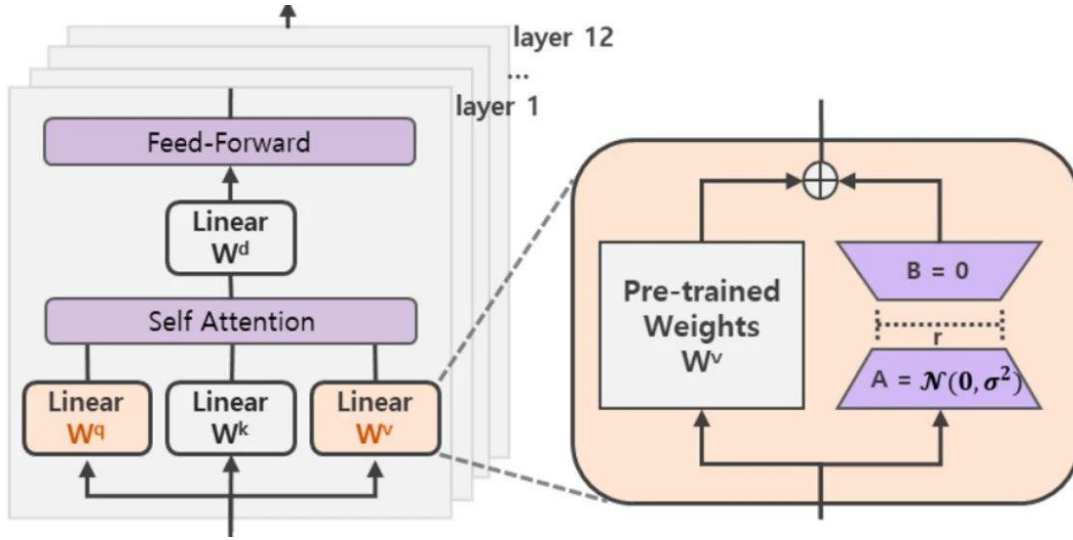


Figure 1: LoRA

Technically, we can insert  $A$  and  $B$  into any layers, such as the query, key, value, output projection layers, or the FFN layers in each attention block. However, the paper suggests that applying LoRA only to query and value projection layers can yield performance comparable to full parameter fine-tuning. The hyperparameter  $r$  typically takes on small values within the range of 4 to 32. By inserting  $A$  and  $B$  into specific layers and keeping the rest of the pre-trained weights frozen, LoRA achieves a substantial reduction in the total number of trainable parameters. For  $r = 8$ , the percentage of parameters that are subject to training accounts for less than 1% of the overall parameters in a typical LM. This remarkable reduction in the number of trainable parameters is a key factor in the efficiency of LoRA as a PEFT technique.

It seems like applying LoRA would add more parameters to the model, thus increasing forward pass costs and inference latency. How does LoRA paper overcome this drawback?

Your answer: In this paper, we introduce a low-rank structure, as  $A$  and  $B$  are low-rank matrices, they introduce significantly fewer parameters than a full matrix. This reduction allows for an increase in parameter count while reducing computational costs. Another method involves updating  $A$  and  $B$ , capitalizing on the parameter efficiency. *LoRA* updates a small fraction of the total number of parameters, which can yield performance comparable to full parameter fine-tuning. By updating the dense matrix  $W_0$ , the multiplication of  $A \times B$  can be optimized during the forward and backward pass. Moreover, *LoRA*'s implementation can leverage efficient linear algebra operations, which are optimized on the hardware, thereby keeping the computational cost low.



### 3.4 Mixed Precision Training

Mixed precision training [?] is a technique used in large neural network training that combines the use of both 16-bit and 32-bit floating-point representations for different parts of the training process. This approach leverages the advantages of lower precision (16-bit) for some computations while still using higher precision (32-bit) where necessary.

As illustrated in Figure ??, during the forward pass, the model converts the FP32 weights to FP16 and computes FP16 activations, and in the backward pass, both activation gradients and weight gradients are calculated in FP16. Subsequently, the FP32 master copy of weights is updated with the FP16 weight gradients. The authors also claim that some operations should be read and written in FP16 but perform arithmetic in FP32 to maintain model accuracy. Read the paper and answer why we need an FP32 master copy of weights and why we need to scale up the loss.

Your answer: An FP32 master copy of weights in mixed precision training is crucial to preserve small gradient values that would otherwise become zero in FP16 due to its limited range and precision. This ensures that all updates, no matter how small, contribute to the learning process. Loss scaling is used to prevent small-magnitude gradients from underflowing to zero in FP16 by temporarily increasing their size, allowing them to be represented in the reduced precision format during back-propagation. Afterward, the gradients are scaled back to their original size before updating the weights.

FP16 operations offer significant speed advantages over FP32 operations, especially on modern GPUs equipped with NVIDIA's Volta and Ampere architectures. These architectures include dedicated hardware, known as Tensor Cores, designed to accelerate FP16 multiplication and accumulation into either FP16 or FP32 outputs. Thus, it is clear that mixed precision training can reduce computation time by utilizing FP16 computation units. In terms of memory consumption, according to the paper:

"Even though maintaining an additional copy of weights increases the memory requirements for the weights by 50% compared with single precision training, impact on overall memory usage is much smaller. For training memory consumption is dominated by activations, due to larger batch sizes and activations of each layer being saved for reuse in the back-propagation pass. Since activations are also stored in half-precision format, the overall memory consumption for training deep neural networks is roughly halved. "

Do you think this statement still holds for LLM fine-tuning?

Your answer: Yes, the statement about mixed precision training reducing overall memory consumption likely still holds for fine-tuning large language models (LLMs). This is because activations, which are stored for use in the backward pass, typically use the most memory during training. Since activations can be stored in FP16 format, this significantly reduces memory usage. Even though maintaining an FP32 copy of the weights increases memory requirements for the weights by 50% compared to FP32 alone, the reduction in memory from using FP16 for activations generally results in a net decrease in overall memory consumption. Additionally, the use of FP16 allows for leveraging specialized GPU hardware like Tensor Cores for faster computations.

## 4 Phase 1: Preliminary Study

Complete all questions in Sections 2 and 3, and upload your answers as a PDF document named `phase_1` to GitHub. It's strongly advised to complete this phase as soon as possible to allow enough time for the subsequent phases.

## 5 Phase 2: LLaMA2 Model Inference

In this phase, we focus on understanding how the LLaMA2 7B model generates text and making specific changes to its code. Our goal is to remove the parts of the code that deal with KV-caching. The LLaMA2 7B model and its tokenizer are located at the directory `project/saifhash_1190/llama2-7b` on the CARC system.

Guidelines:

1. Understanding the Code Base: Familiarize yourself with the model's architecture, the tokenization process, batch generation, and the decoding mechanism. Utilize debugging tools to inspect values and



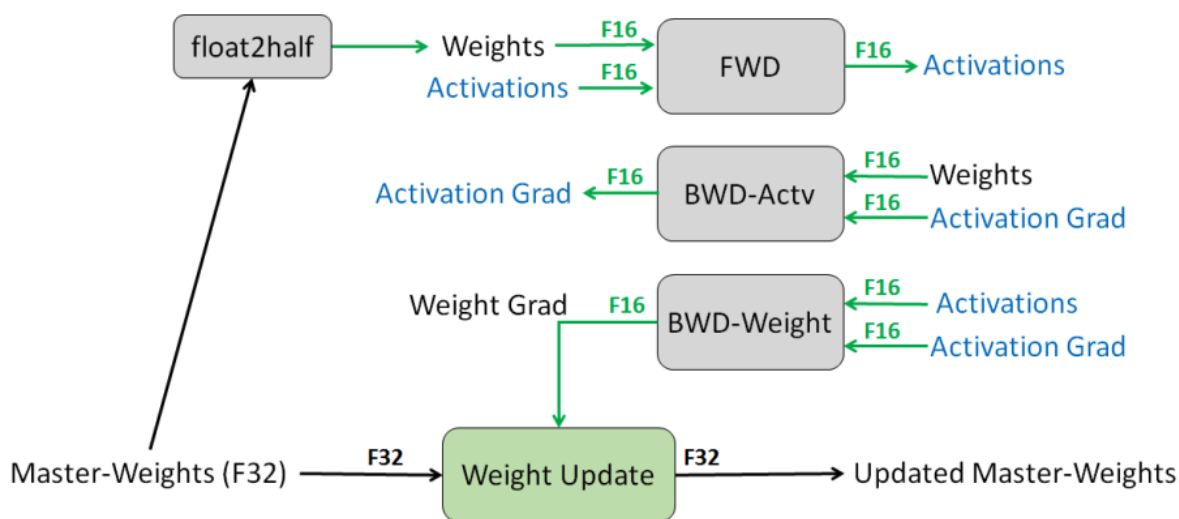


Figure 2: Mixed precision training

flow as necessary.

2. Identifying KV-Caching Features: Look through the code to locate all components and functions related to KV-caching.
3. Modifying the Code: Carefully remove the KV-caching features from the code. Make sure the model can still generate text without these features.
4. Testing: After making changes, test the model with sample prompts to ensure it still works correctly.

Deliverables: Submit a PDF document named `phase_2` containing the details of the changes you've made to the code, along with the prompts and the outputs from testing the generation process.

## 6 Phase 3: LLaMA2 Model Training

In this phase, your task is to conduct instruction tuning on the LLaMA2 7B model using the (partial) Alpaca dataset, which enhances and expands the capabilities of the LLaMA2 model. By applying all previously discussed optimization techniques, fine-tuning LLMs will become feasible, even on a single GPU. We will avoid using high-level libraries such as HuggingFace, which abstracts away many important details. Instead, we will solely rely on PyTorch's built-in functions. This approach will ensure you gain a comprehensive understanding of the process involved in the instruction tuning of LLMs.

Guidelines:

1. Initial Setup: Begin by instantiating only 1 decoder layer of the LLaMA2 model. Test all your code on a less powerful yet more accessible GPU, the P100 on CARC. This approach prevents out-of-memory errors even before applying optimization techniques. Once you've ensured an end-to-end training process that is bug-free and incorporates all four optimization techniques, you can scale up to the full-size LLaMA2 model, which includes 32 decoder layers, and utilize a more powerful GPU, the A100, with 40GB of RAM, on CARC system.
2. End-to-End Instruction Tuning Flow: Create a file named `finetuning.py`. Implement the end-to-end instruction tuning workflow in PyTorch. For data preprocessing and the creation of supervised data loaders, refer to the official Alpaca repo.

3. Training Iteration Loop: Replace the HuggingFace `trainer` object used in the Alpaca repo with your own implementation from scratch. To compute loss using PyTorch functions, refer to the HuggingFace’s implementation.
4. Gradient Accumulation and Mixed Precision Training: Implement gradient accumulation and mixed precision training using PyTorch’s AMP package. Refer to the AMP Recipe and AMP Examples for guidance.
5. LoRA Linear Layer Module: Implement the LoRA linear layer module by referring to the official implementation. Create a file named `lora.py` under the `llama/model` path. Convert your model into a PEFT model by replacing the Q and V projection layers in the LLaMA2 model with your LoRA `Linear` modules. Freeze all model parameters except for `LoRA_A` and `LoRA_B` and report the percentage of trainable parameters.
6. Gradient Checkpointing: Utilize PyTorch’s `torch.utils.checkpoint` API for inserting checkpoints. Decide which layer(s) to apply checkpointing to and describe your approach in the report. Refer to this tutorial for more information.
7. Model Fine-Tuning: Apply all the aforementioned techniques to your LLaMA2 7B model and fine-tune it on the Alpaca Dataset using the A100 GPU. To manage time constraints for educational purposes, extract the first 200 samples from the dataset as your training set.
8. Hyperparameters: Set  $learning\_rate = 1e-5$ ,  $batch\_size = 1$  and  $gradient\_accumulation\_step = 8$ . For LoRA configuration, set  $r = 16$ ,  $alpha = 32$ , and  $dropout\_rate = 0.05$ .

Deliverables: Submit a PDF file named `phase_3` that includes the implementation details you have completed, along with the prompt and the result used to test the fine-tuned model. In the same PDF file, provide the following analysis and measurements. For Table ??, fill in the blanks with  $\uparrow$  to signify an increase in resource usage,  $\downarrow$  for a decrease, or  $-$  to indicate no change. For Table ??, measure and report the peak memory usage in MB and the runtime in seconds required to process only 200 training samples on the A100 GPU across each combination of techniques applied in the instruction tuning of the LLaMA2 7B model. If an out-of-memory error happens, then simply mark  $\times$  in the blanks. You can turn off gradient checkpointing for this step.

		Grad. Accumulation	Grad. Checkpoint	Mixed Precision	LoRA
Memory	parameter				
	activation				
	gradient				
	optimizer state				
Computation					

Table 1: System performance analysis

GA	OFF				ON			
MP	OFF		ON		OFF		ON	
LoRA	OFF	ON	OFF	ON	OFF	ON	OFF	ON
Peak Mem								
Runtime								

Table 2: System performance measurement